

EC2 AUTO SCALING AND LOAD BALANCING

CLOUD PROJECT REPORT

Presented by: Vanshika Munjal & Vaani Jindal (Group 37)

500121784

500119144

Batch: B2 CCVT

INTRODUCTION

Cloud Computing Overview:

Cloud computing is the on-demand delivery of IT resources over the Internet with pay-as-you-go pricing. Instead of buying, owning, and maintaining physical data centers and servers, organizations can access computing services—such as servers, storage, databases, networking, software, and analytics—from a cloud provider.

Amazon Web Services (AWS) Overview:

Amazon Web Services (AWS) is the world's most comprehensive and widely adopted cloud platform. It offers over 200 fully featured services, including computing, storage, networking, machine learning, security, and analytics.

PURPOSE OF PROJECT

This project focuses on building a scalable and highly available system using AWS services. By implementing EC2 Auto Scaling, we ensure that our application automatically adjusts the number of instances based on demand. Load Balancing, on the other hand, distributes incoming traffic efficiently across multiple EC2 instances, improving performance and reliability. Together, these features help build a fault-tolerant and cost-effective cloud infrastructure.

PROJECT GOALS

The goal of this project is to design and deploy a scalable, resilient, and cost-efficient cloud infrastructure using AWS EC2 Auto Scaling and Elastic Load Balancing. This system is intended to handle unpredictable traffic patterns by automatically adjusting compute capacity, ensuring high availability, fault tolerance, and optimal performance at all times.

Through this project, we aim to demonstrate how cloud-native technologies can eliminate manual resource management, reduce downtime, and support real-world applications with dynamic workloads effectively.

PROJECT OBJECTIVES

- Implement EC2 Auto Scaling to automatically add or remove instances based on demand.
- Integrate IAM, S3, EBS, and Lambda
- Configure an Elastic Load Balancer (ELB) to distribute incoming traffic evenly across instances.
- Optimize resource usage and cost by scaling the infrastructure dynamically.
- Monitor system performance using AWS tools like CloudWatch for real-time insights.
- Ensure best security practices using IAM roles and encryption

AWS SERVICES USED

IAM (Identity and Access Management)

Roles & Permissions setup:

To ensure secure access to AWS services, Identity and Access Management (IAM) was configured as follows:

- Created IAM Group: EC2Group with AmazonEC2FullAccess policy.
- Created IAM User: Assigned to EC2Group, enabling programmatic access with generated Access Key and Secret Key.
- IAM Role for EC2 (EC2AutoScalingRole):
- Policy Attached: AmazonEC2RoleforSSM and AmazonS3FullAccess.

This role was associated with the EC2 Launch Template to allow EC2 instances to access S3 and use SSM for remote access.

EC2 Instance Configurations:

Amazon EC2 was used to launch virtual servers with the following setup:

- Launch Template: Named MyAutoScaleTemplate with:
- AMI: Amazon Linux 2
- Instance type: t2.micro
- Key pair: Used for SSH access
- Security Group: Allows SSH (port 22) and HTTP (port 80)
- IAM instance profile: EC2AutoScalingRole
- Auto Scaling Group:
- Min Size: 1, Max Size: 5, Desired Capacity: 2
- Subnets used: subnet-09aa416d352d8792e, subnet-081ae8c6100435c9a, subnet-0e68c9beb9282a1ff
- Scaling policy triggered by CPU utilization

AWS SERVICES USED

S3 Bucket Settings and EBS Attachment:

S3 Bucket:

- Created using CLI
- Used for file storage and syncing files with EC2 instances (aws s3 sync, cp, rm commands)
- Versioning and lifecycle rules were considered for optimization.

EBS Volume:

- Created and attached to EC2 instances
- Mounted to /data directory after formatting with mkfs
- Commands used: aws ec2 create-volume, attach-volume, followed by SSH-based Linux commands to mount and persist.

Lambda Function:

- Function Name: EC2CleanupFunction
- Purpose: Automatically stops/terminates idle EC2 instances based on CloudWatch Events
- Runtime: Python 3.x
- Trigger: Scheduled EventBridge rule EC2CleanupRule (e.g., every 1 hour)

Permissions:

- Lambda permission added for EventBridge to invoke it (lambda:InvokeFunction)
- IAM role with ec2:DescribeInstances, ec2:StopInstances, and logs:/* permissions

IMPLEMENTATION DETAILS AND CODE SNIPPETS

Configuring AWS

```
C:\Users\VAANI JINDAL>aws configure
AWS Access Key ID [*****VKVW]: AKIA5G2VG5EDEAD2FZ6P
AWS Secret Access Key [*****QQVg]: hzV9FOHzr6Nrd14r/VfAIIVzbcdUaRScuKnzY3oBf
Default region name [ap-south-1]: ap-south-1
Default output format [json]: json
```

```
C:\Users\VAANI JINDAL>aws sts get-caller-identity
{
    "UserId": "AIDA5G2VG5EDLGUMCLZ7I",
    "Account": "908027422982",
    "Arn": "arn:aws:iam::908027422982:user/user_1"
}
```

The command `aws sts get-caller-identity` is used to retrieve details about the AWS identity making the request. It helps you confirm which AWS account, user, or role is being used when executing AWS CLI commands.

```
C:\Users\VAANI JINDAL>notepad trust-policy.json  
C:\Users\VAANI JINDAL>type trust-policy.json  
{  
    "Version": "2012-10-17",  
    "Statement": [  
        {  
            "Effect": "Allow",  
            "Principal": {  
                "Service": "ec2.amazonaws.com"  
            },  
            "Action": "sts:AssumeRole"  
        }  
    ]  
}
```

Setting up IAM Role

The trust-policy.json file is created to define the Trust Policy for an IAM Role in AWS.

- notepad trust-policy.json created json file
- type trust-policy.jsonconfirmation that created successfully

```
C:\Users\VAANI JINDAL>aws iam create-role --role-name EC2AutoScalingRole --assume-role-policy-document file:///"C:\\\\Users\\\\VAANI JINDAL\\\\OneDrive\\\\Desktop\\\\trust-policy.json"  
{  
    "Role": {  
        "Path": "/",  
        "RoleName": "EC2AutoScalingRole",  
        "RoleId": "AROA5G2VG5EDPAZTDQ7YK",  
        "Arn": "arn:aws:iam::908027422982:role/EC2AutoScalingRole",  
        "CreateDate": "2025-04-03T14:55:06+00:00",  
        "AssumeRolePolicyDocument": {  
            "Version": "2012-10-17",  
            "Statement": [  
                {  
                    "Effect": "Allow",  
                    "Principal": {  
                        "Service": "ec2.amazonaws.com"  
                    },  
                    "Action": "sts:AssumeRole"  
                }  
            ]  
        }  
    }  
}
```

IAM user_1

The screenshot shows the AWS Identity and Access Management (IAM) service interface. The top navigation bar includes the AWS logo, a search bar with placeholder 'Search' and keyboard shortcut '[Alt+S]', and various global settings like 'Global' and 'VAANI JINDAL'. The main navigation on the left is under 'Identity and Access Management (IAM)' and includes 'Dashboard', 'Access management' (which is expanded to show 'User groups', 'Users', 'Roles', and 'Policies'), and 'Policies'. The current view is for a specific user named 'user_1', which is selected in the breadcrumb trail ('IAM > Users > user_1'). The main content area displays the 'user_1' summary information. The ARN is listed as 'arn:aws:iam::908027422982:user/user_1'. The user was created on 'April 03, 2025, 17:40 (UTC+05:30)'. Console access is 'Disabled'. There is one security credential listed: 'Access key 1' (AKIA5G2VG5EDEAD2FZ6P - Active, used today, 20 hours old). Other tabs at the bottom include 'Permissions', 'Groups', 'Tags (1)', 'Security credentials', and 'Last Accessed'.

Identity and Access Management (IAM)

user_1 Info

Summary

ARN
arn:aws:iam::908027422982:user/user_1

Created
April 03, 2025, 17:40 (UTC+05:30)

Console access
Disabled

Last console sign-in
-

Access key 1
AKIA5G2VG5EDEAD2FZ6P - Active
Used today. 20 hours old.

Access key 2
Create access key

Permissions Groups Tags (1) Security credentials Last Accessed

Attaching Required Policies to IAM Role

```
C:\Users\VAANI JINDAL>aws iam attach-role-policy --role-name EC2AutoScalingRole --policy-arn arn:aws:iam::aws:policy/AmazonEC2FullAccess

C:\Users\VAANI JINDAL>aws iam attach-role-policy --role-name EC2AutoScalingRole --policy-arn arn:aws:iam::aws:policy/service-role/AWSLambdaBasicExecutionRole

C:\Users\VAANI JINDAL>aws iam attach-role-policy --role-name EC2AutoScalingRole --policy-arn arn:aws:iam::aws:policy/AmazonS3FullAccess
```

Verification of the attached policies

```
C:\Users\VAANI JINDAL>aws iam list-attached-role-policies --role-name EC2AutoScalingRole
{
    "AttachedPolicies": [
        {
            "PolicyName": "AmazonEC2FullAccess",
            "PolicyArn": "arn:aws:iam::aws:policy/AmazonEC2FullAccess"
        },
        {
            "PolicyName": "AWSLambdaBasicExecutionRole",
            "PolicyArn": "arn:aws:iam::aws:policy/service-role/AWSLambdaBasicExecutionRole"
        },
        {
            "PolicyName": "AmazonS3FullAccess",
            "PolicyArn": "arn:aws:iam::aws:policy/AmazonS3FullAccess"
        }
    ]
}
```

Launch EC2 Instance Process

```
C:\Users\VAANI JINDAL> aws ec2 create-key-pair --key-name MyEC2Key --query "KeyMaterial" --output text > MyEC2Key.pem

C:\Users\VAANI JINDAL>chmod 400 MyEC2Key.pem
'chmod' is not recognized as an internal or external command,
operable program or batch file.

C:\Users\VAANI JINDAL>icacls MyEC2Key.pem /inheritance:r
processed file: MyEC2Key.pem
Successfully processed 1 files; Failed processing 0 files

C:\Users\VAANI JINDAL>icacls MyEC2Key.pem /grant:r "%username%:F"
processed file: MyEC2Key.pem
Successfully processed 1 files; Failed processing 0 files

C:\Users\VAANI JINDAL>icacls MyEC2Key.pem /deny "NT AUTHORITY\SYSTEM:(F)"
processed file: MyEC2Key.pem
Successfully processed 1 files; Failed processing 0 files
```

- Setting up key-pair
- Remove inherited permissions (so only your user can access it)
- Grant full access only to your current user
- Deny access to the SYSTEM account (extra security)

Now the .pem file is secured and we can use it to connect to ourEC2 instance!

```
C:\Users\VAANI JINDAL>dir MyEC2Key.pem
Volume in drive C is Windows
Volume Serial Number is 3E9D-5BFC

Directory of C:\Users\VAANI JINDAL

03-04-2025  23:38           1,702 MyEC2Key.pem
              1 File(s)      1,702 bytes
              0 Dir(s)  13,773,385,728 bytes free
```

Running this command in cmd to verify if .pem file exists

Created a Security Group

A Security Group is required to allow inbound and outbound traffic to your EC2 instance. This creates a security group named MySecurityGroup.

```
C:\Users\VAANI JINDAL>aws ec2 create-security-group --group-name MySecurityGroup --description "Security group for EC2"
{
    "GroupId": "sg-0f4320053ede4b289",
    "SecurityGroupArn": "arn:aws:ec2:ap-south-1:908027422982:security-group/sg-0f4320053ede4b289"
}
```

```
C:\Users\VAANI JINDAL>aws ec2 authorize-security-group-ingress --group-name MySecurityGroup --protocol tcp --port 22  
--cidr 0.0.0.0/0  
{  
    "Return": true,  
    "SecurityGroupRules": [  
        {  
            "SecurityGroupRuleId": "sgr-044ffd45a47f6167e",  
            "GroupId": "sg-0f4320053ede4b289",  
            "GroupOwnerId": "908027422982",  
            "IsEgress": false,  
            "IpProtocol": "tcp",  
            "FromPort": 22,  
            "ToPort": 22,  
            "CidrIpv4": "0.0.0.0/0",  
            "SecurityGroupRuleArn": "arn:aws:ec2:ap-south-1:908027422982:security-group-rule/sgr-044ffd45a47f6167e"  
        }  
    ]  
}  
  
C:\Users\VAANI JINDAL>aws ec2 authorize-security-group-ingress --group-name MySecurityGroup --protocol tcp --port 80  
--cidr 0.0.0.0/0  
{  
    "Return": true,  
    "SecurityGroupRules": [  
        {  
            "SecurityGroupRuleId": "sar-032b21b7aad4aa1b3"  
        }  
    ]  
}
```

This allows you to connect via SSH from any IP.
Now, your EC2 can be accessed via SSH and HTTP

To get latest amazon Linux 2 AMI Id

```
C:\Users\VAANI JINDAL>aws ec2 describe-images --owners amazon --filters "Name=name,Values=amzn2-ami-hvm-*--x86_64-gp2"  
--region ap-south-1 --query "Images | sort_by(@, &CreationDate)[-1].ImageId" --output text  
ami-0019610b50e1fe6bb
```

Launch EC2 Instance (region ap-south-1)

```
C:\Users\VAANI JINDAL>aws ec2 run-instances --image-id ami-0019610b50e1fe6bb --instance-type t2.micro --key-name MyEC2Key --security-group-ids sg-0f4320053ede4b289 --count 1 --region ap-south-1
{
    "ReservationId": "r-034bca6995523fe53",
    "OwnerId": "908027422982",
    "Groups": [],
    "Instances": [
        {
            "Architecture": "x86_64",
            "BlockDeviceMappings": [],
            "ClientToken": "d2fc96a0-b0e6-4005-88cc-831ed17fa06f",
            "EbsOptimized": false,
            "EnaSupport": true,
            "Hypervisor": "xen",
            "NetworkInterfaces": [
                {
                    "Attachment": {
                        "AttachTime": "2025-04-03T18:50:02+00:00",
                        "AttachmentId": "eni-attach-0e7098bleecf369cf",
                        "DeleteOnTermination": true,
                        "DeviceIndex": 0,
                        "Status": "attaching",
                        "NetworkCardIndex": 0
                    },
                    "Description": "",
                    "Groups": [

```

Connect to Your EC2 Instance

```
C:\Users\VAANI JINDAL>aws ec2 describe-instances --query "Reservations[].Instances[][InstanceId,PublicIpAddress]" --output table
-----|-----|-----|-----|
|     DescribeInstances      | |-----|-----|
+-----+-----+-----+-----+
| i-069bf6155cdbe2690 | None   | |-----|-----|
| i-006e7bd91fc2e574b | 3.110.223.227 | |-----|-----|
+-----+-----+-----+
```

EC2 Instance

The screenshot shows the AWS EC2 Instance Details page for an instance with ID **i-006e7bd91fc2e574b**. The instance is currently **Running** (indicated by a green checkmark). Key details include:

- Public IPv4 address:** 3.110.223.227 | [open address](#)
- Private IP4 addresses:** 172.31.5.120
- Public IPv4 DNS:** ec2-3-110-223-227.ap-south-1.compute.amazonaws.com | [open address](#)
- Hostname type:** IP name: ip-172-31-5-120.ap-south-1.compute.internal
- Answer private resource DNS name:** -
- Private IP DNS name (IPv4 only):** ip-172-31-5-120.ap-south-1.compute.internal
- Instance type:** t2.micro

The left sidebar shows the navigation path: EC2 > Instances > i-006e7bd91fc2e574b. The Instances section is expanded, showing options like Instances, Instance Types, Launch Templates, etc.

Verify public ip

```
C:\Users\VAANI JINDAL>aws ec2 describe-instances --instance-ids i-006e7bd91fc2e574b --query "Reservations[].Instances[].[State.Name, PublicIpAddress]" --output table
-----
|   DescribeInstances   |
+-----+-----+
| running | 3.110.223.227 |
+-----+-----+
```

To connect

```
C:\Users\VAANI JINDAL>ssh -i "MyEC2Key.pem" ec2-user@3.110.223.227
The authenticity of host '3.110.223.227 (3.110.223.227)' can't be established.
ED25519 key fingerprint is SHA256:K7tUsUqKOF04np937kwBHTSHzWlq0YuSM4PBseBZZrw.
This key is not known by any other names.
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added '3.110.223.227' (ED25519) to the list of known hosts.
```

AWS configure and then creating EBS volume

```
[ec2-user@ip-172-31-5-120 ~]$ aws configure
AWS Access Key ID [None]: AKIA5G2VG5EDEAD2FZ6P
AWS Secret Access Key [None]: hzV9FOHzr6Nrd14r/VfAIIVzbcbgUaRScuKnzY3oBf
Default region name [None]: ap-south-1
Default output format [None]: json
[ec2-user@ip-172-31-5-120 ~]$ aws ec2 create-volume --availability-zone ap-south-1b --size 10 --volume-type gp2 --region ap-south-1
{
    "AvailabilityZone": "ap-south-1b",
    "MultiAttachEnabled": false,
    "Tags": [],
    "Encrypted": false,
    "VolumeType": "gp2",
    "VolumeId": "vol-0b286b8328c3fae22",
    "State": "creating",
    "Iops": 100,
    "SnapshotId": "",
    "CreateTime": "2025-04-03T19:30:42.000Z",
    "Size": 10
}
```

Attaching the EBS Volume

```
[ec2-user@ip-172-31-5-120 ~]$ aws ec2 attach-volume --volume-id vol-0b286b8328c3fae22 --instance-id i-006e7bd91fc2e574b --device /dev/xvdf --region ap-south-1
{
    "AttachTime": "2025-04-03T19:32:45.219Z",
    "InstanceId": "i-006e7bd91fc2e574b",
    "VolumeId": "vol-0b286b8328c3fae22",
    "State": "attaching",
    "Device": "/dev/xvdf"
}
```

EBS volume

The screenshot shows the AWS EC2 Volumes page for an EBS volume named **vol-0a2331ad29fc852c0**. The volume is currently **In-use** and was created on **Fri Apr 04 2025 00:20:03 GMT+0530 (India Standard Time)**. It has a size of **8 GiB** and an IOPS of **100**, and is attached to the **i-006e7bd91fc2e574b** instance at **/dev/xvda**. The volume is of type **gp2** and is located in the **ap-south-1b** availability zone. The status check is **Okay**.

Details

Volume ID	Size	Type	Status check
vol-0a2331ad29fc852c0	8 GiB	gp2	Okay
AWS Compute Optimizer finding	Volume state	IOPS	Throughput
Opt-in to AWS Compute Optimizer for recommendations. Learn more	In-use	100	-
Fast snapshot restored	Availability Zone	Created	Multi-Attach enabled
No	ap-south-1b	Fri Apr 04 2025 00:20:03 GMT+0530 (India Standard Time)	No
Attached resources	Outposts ARN	Managed	Operator
i-006e7bd91fc2e574b: /dev/xvda (attached)	-	false	-

To enable automatic backup

```
[ec2-user@ip-172-31-5-120 ~]$ aws ec2 create-snapshot --volume-id vol-0b286b8328c3fae22 --description "Automatic backup"
{
    "Description": "Automatic backup",
    "Tags": [],
    "Encrypted": false,
    "VolumeId": "vol-0b286b8328c3fae22",
    "State": "pending",
    "VolumeSize": 10,
    "StartTime": "2025-04-03T20:12:25.361Z",
    "Progress": "",
    "OwnerId": "908027422982",
    "SnapshotId": "snap-0951794db846f927a"
}
```

To schedule backups automatically

```
[ec2-user@ip-172-31-5-120 ~]$ aws backup create-backup-vault --backup-vault-name MyBackupVault
{
    "BackupVaultArn": "arn:aws:backup:ap-south-1:908027422982:backup-vault:MyBackupVault",
    "CreationDate": 1743711954.947,
    "BackupVaultName": "MyBackupVault"
}
[ec2-user@ip-172-31-5-120 ~]$ aws backup create-backup-plan --backup-plan '{"BackupPlanName": "MyBackupPlan", "Rules": [{"RuleName": "DailyBackup", "TargetBackupVaultName": "MyBackupVault", "ScheduleExpression": "cron(0 12 * * ? *)", "StartWindowMinutes": 60, "CompletionWindowMinutes": 120}]}'
{
    "BackupPlanArn": "arn:aws:backup:ap-south-1:908027422982:backup-plan:e4086bd7-482e-43f5-9e08-41e5a773cba4",
    "VersionId": "M2JlOGI1ZDktNzEzZC00ZDQ5LWE0YzQtNzhlyjJkM2FiZmYz",
    "CreationDate": 1743712036.105,
    "BackupPlanId": "e4086bd7-482e-43f5-9e08-41e5a773cba4"
```

These permissions attached to iam user

- AWSBackupFullAccess
- AWSBackupOperatorAccess

To schedule backups automatically:
AWS Backup stores all backups in a Backup Vault. By default, AWS creates a vault named "Default", but in our case, it might not exist.

command 1 for Create a Backup Vault
command 2 for To schedule backups:

CREATING S3 BUCKET

```
. Please try again.  
[ec2-user@ip-172-31-5-120 ~]$ aws s3 mb s3://autoscaling-bucket-new1 --region ap-south-1  
make_bucket: autoscaling-bucket-new1
```

Cd command

```
C:\Users\VAANI JINDAL\OneDrive\Desktop>cd C:\Users\VAANI JINDAL\OneDrive\Desktop
```

Upload a File to S3 using Python

```
C:\Users\VAANI JINDAL\OneDrive\Desktop>python upload_s3.py  
File C:\Users\VAANI JINDAL\OneDrive\Desktop\testfile.txt uploaded to S3 bucket autoscaling-bu  
cket-new1 as test-file.txt
```

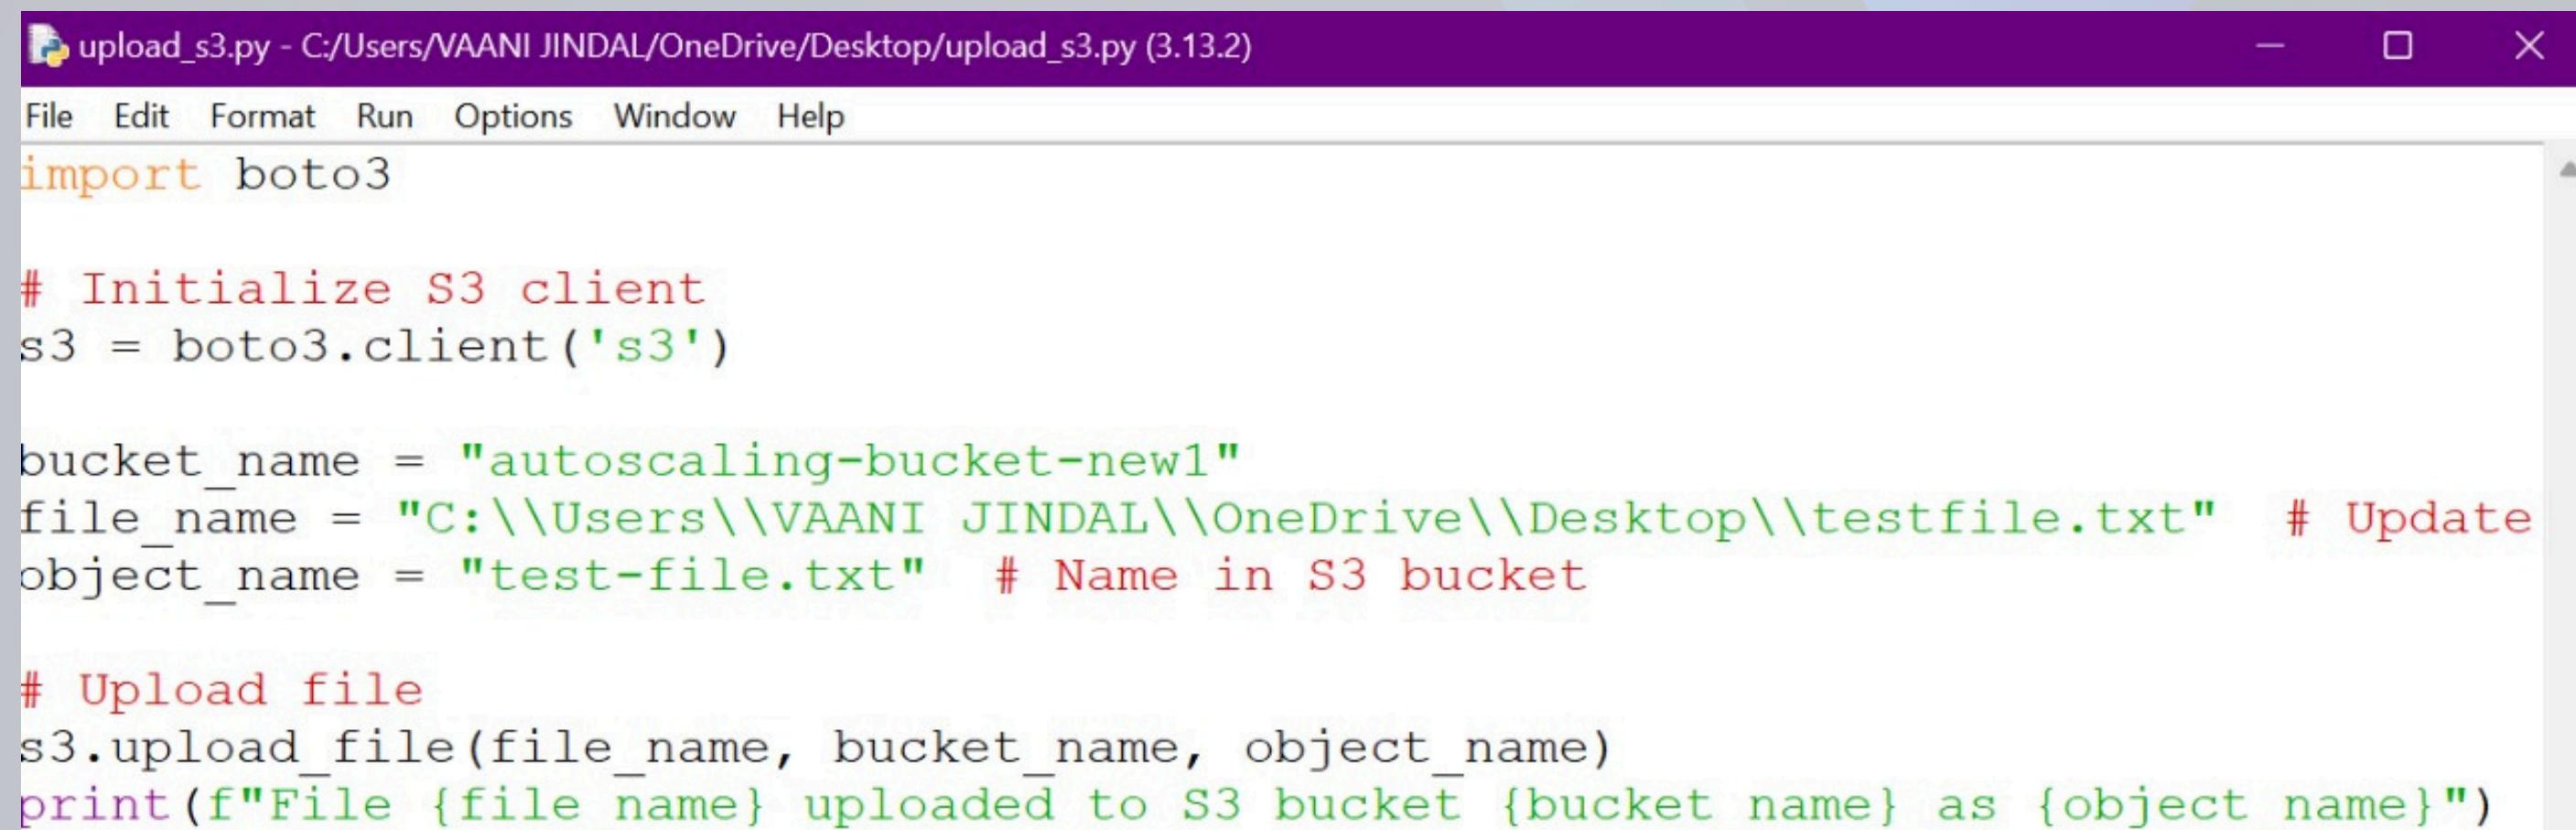
S3 BUCKET

The screenshot shows the AWS S3 console interface. At the top, there's a navigation bar with the AWS logo, a search bar, and various icons. Below the navigation bar, the path is shown as Amazon S3 > Buckets > autoscaling-bucket-new1. The main area is titled "Objects (1)" and displays a single file named "test-file.txt". The file details are as follows:

Name	Type	Last modified	Size	Storage class
test-file.txt	txt	April 4, 2025, 09:44:28 (UTC+05:30)	24.0 B	Standard

On the left sidebar, under "Amazon S3", the "General purpose buckets" section is expanded, showing options like Directory buckets, Table buckets, Access Grants, Access Points, Object Lambda Access Points, Multi-Region Access Points, Batch Operations, and IAM Access Analyzer for S3.

Upload a File to S3 using Python



The screenshot shows a Python code editor window with the title "upload_s3.py - C:/Users/VAANI JINDAL/OneDrive/Desktop/upload_s3.py (3.13.2)". The menu bar includes File, Edit, Format, Run, Options, Window, and Help. The code itself uses the Boto3 library to upload a file from a local path to an S3 bucket. It defines variables for the bucket name, file name, and object name, then uses the `s3.upload_file` method to perform the upload and prints a confirmation message.

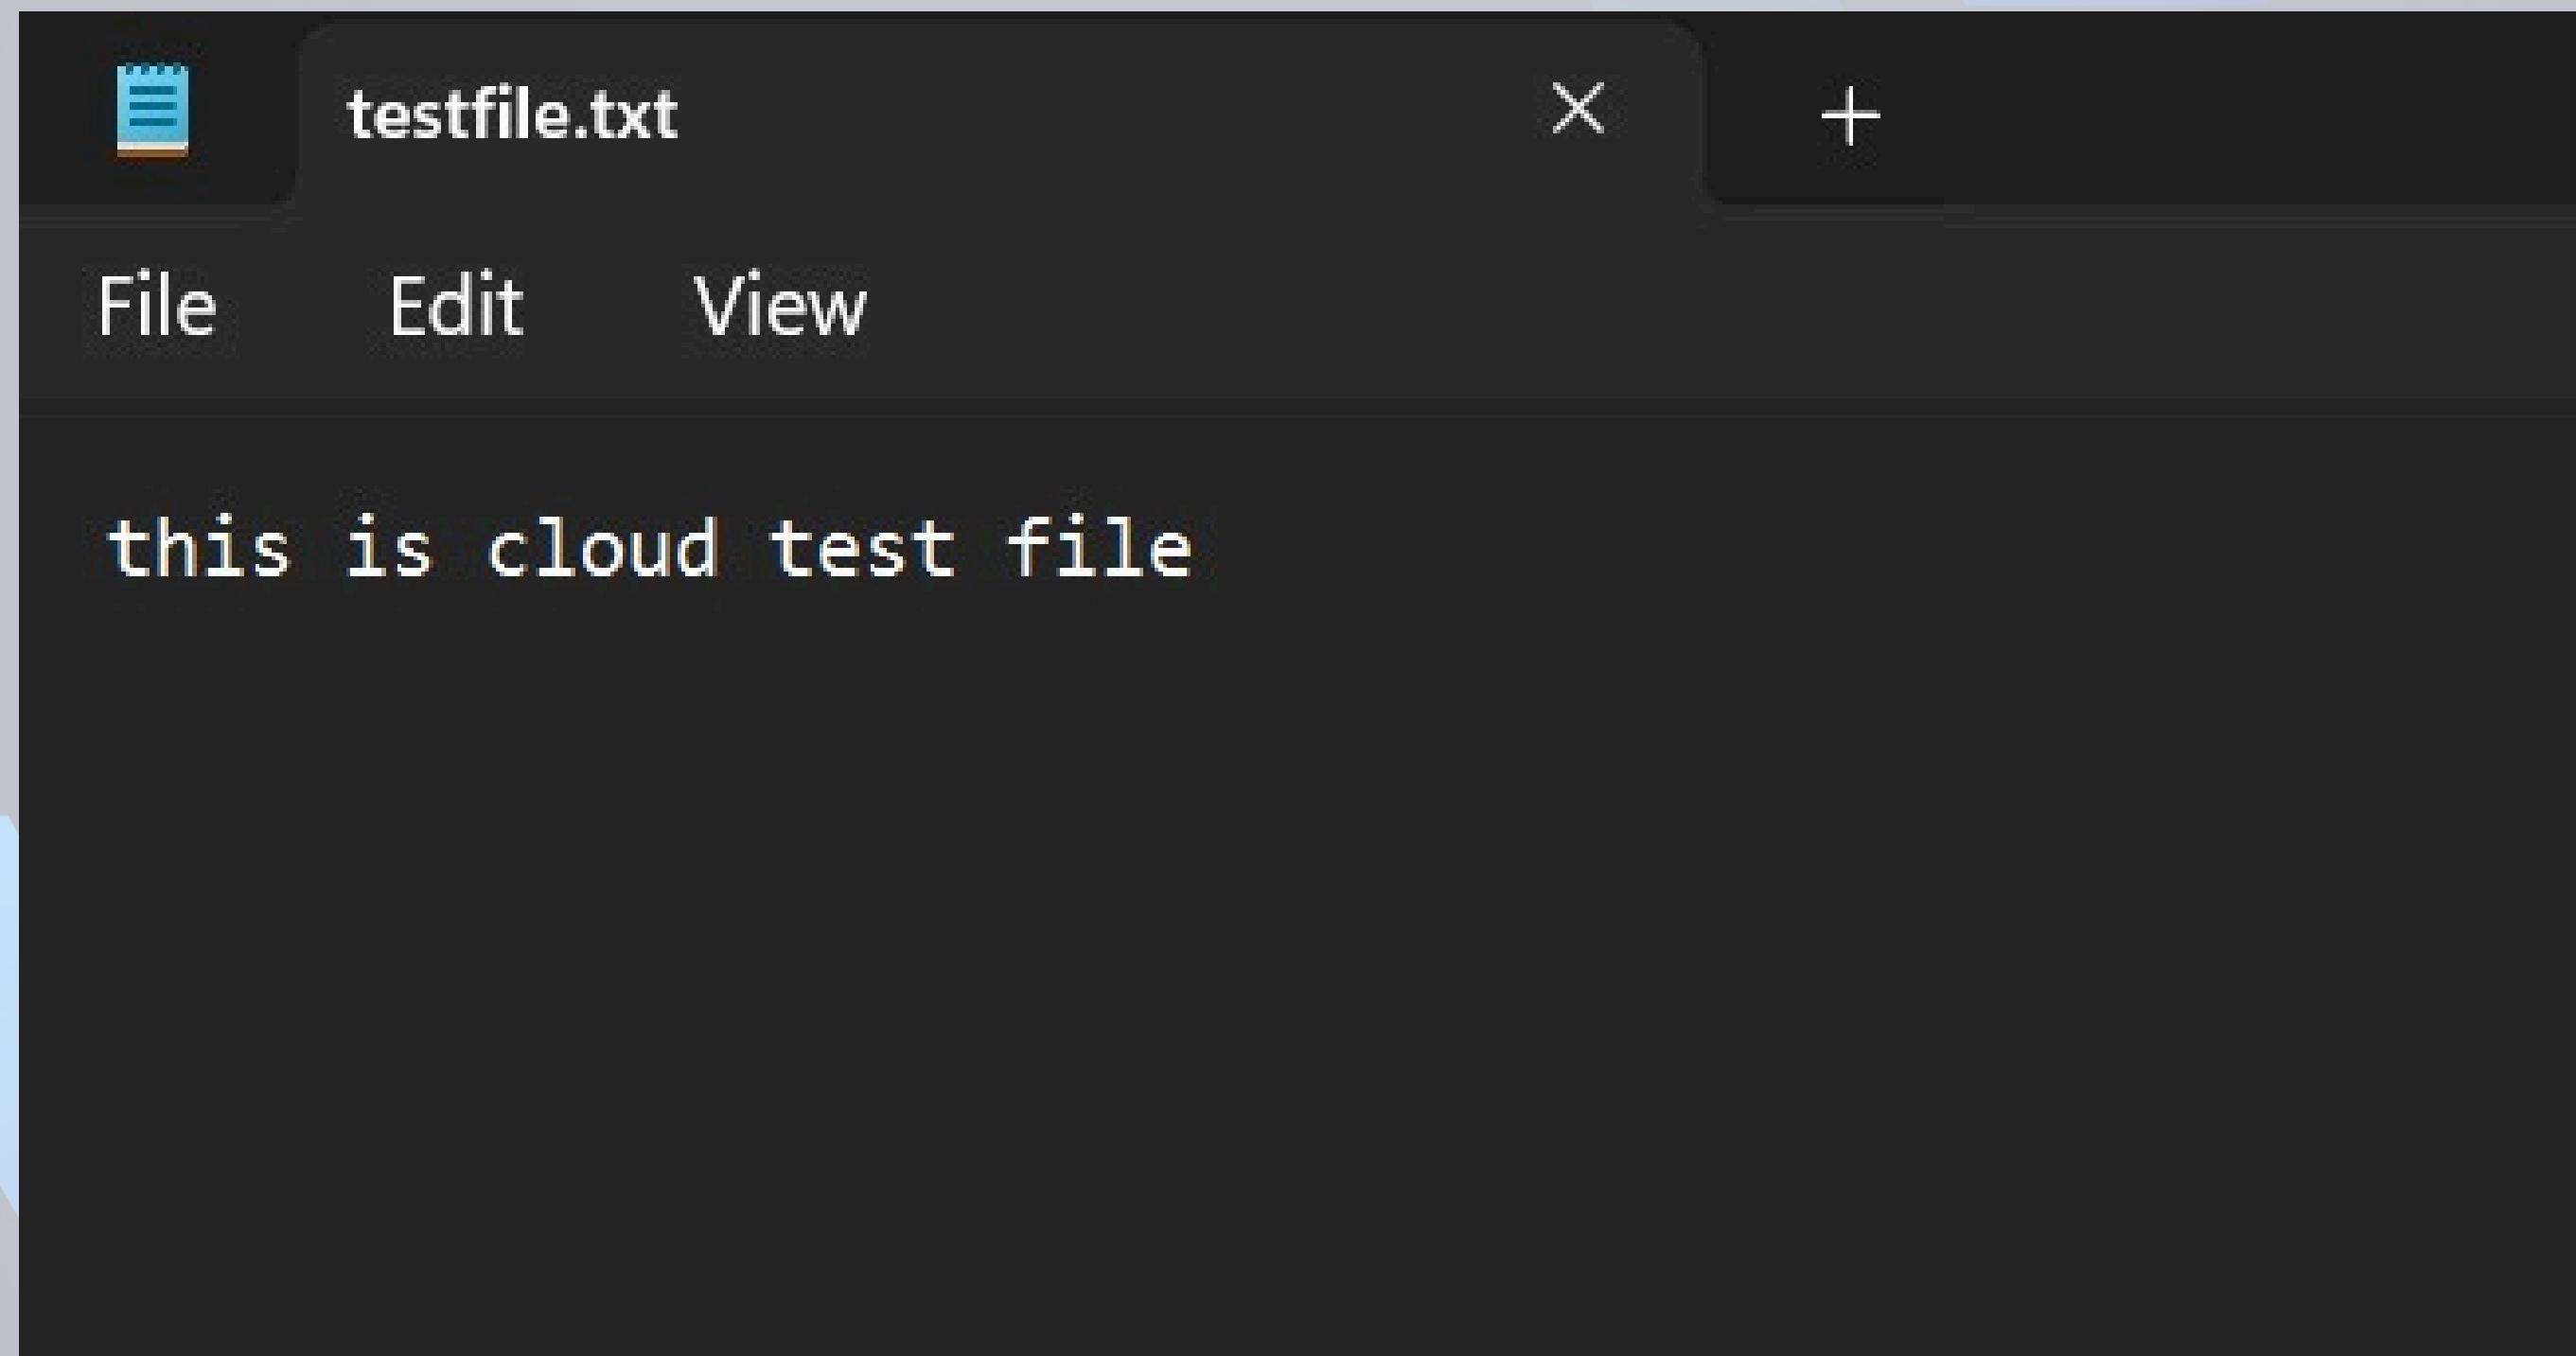
```
import boto3

# Initialize S3 client
s3 = boto3.client('s3')

bucket_name = "autoscaling-bucket-new1"
file_name = "C:\\\\Users\\\\VAANI JINDAL\\\\OneDrive\\\\Desktop\\\\testfile.txt" # Update
object_name = "test-file.txt" # Name in S3 bucket

# Upload file
s3.upload_file(file_name, bucket_name, object_name)
print(f"File {file_name} uploaded to S3 bucket {bucket_name} as {object_name}")
```

File uploaded to S3 Bucket (testfile.txt)

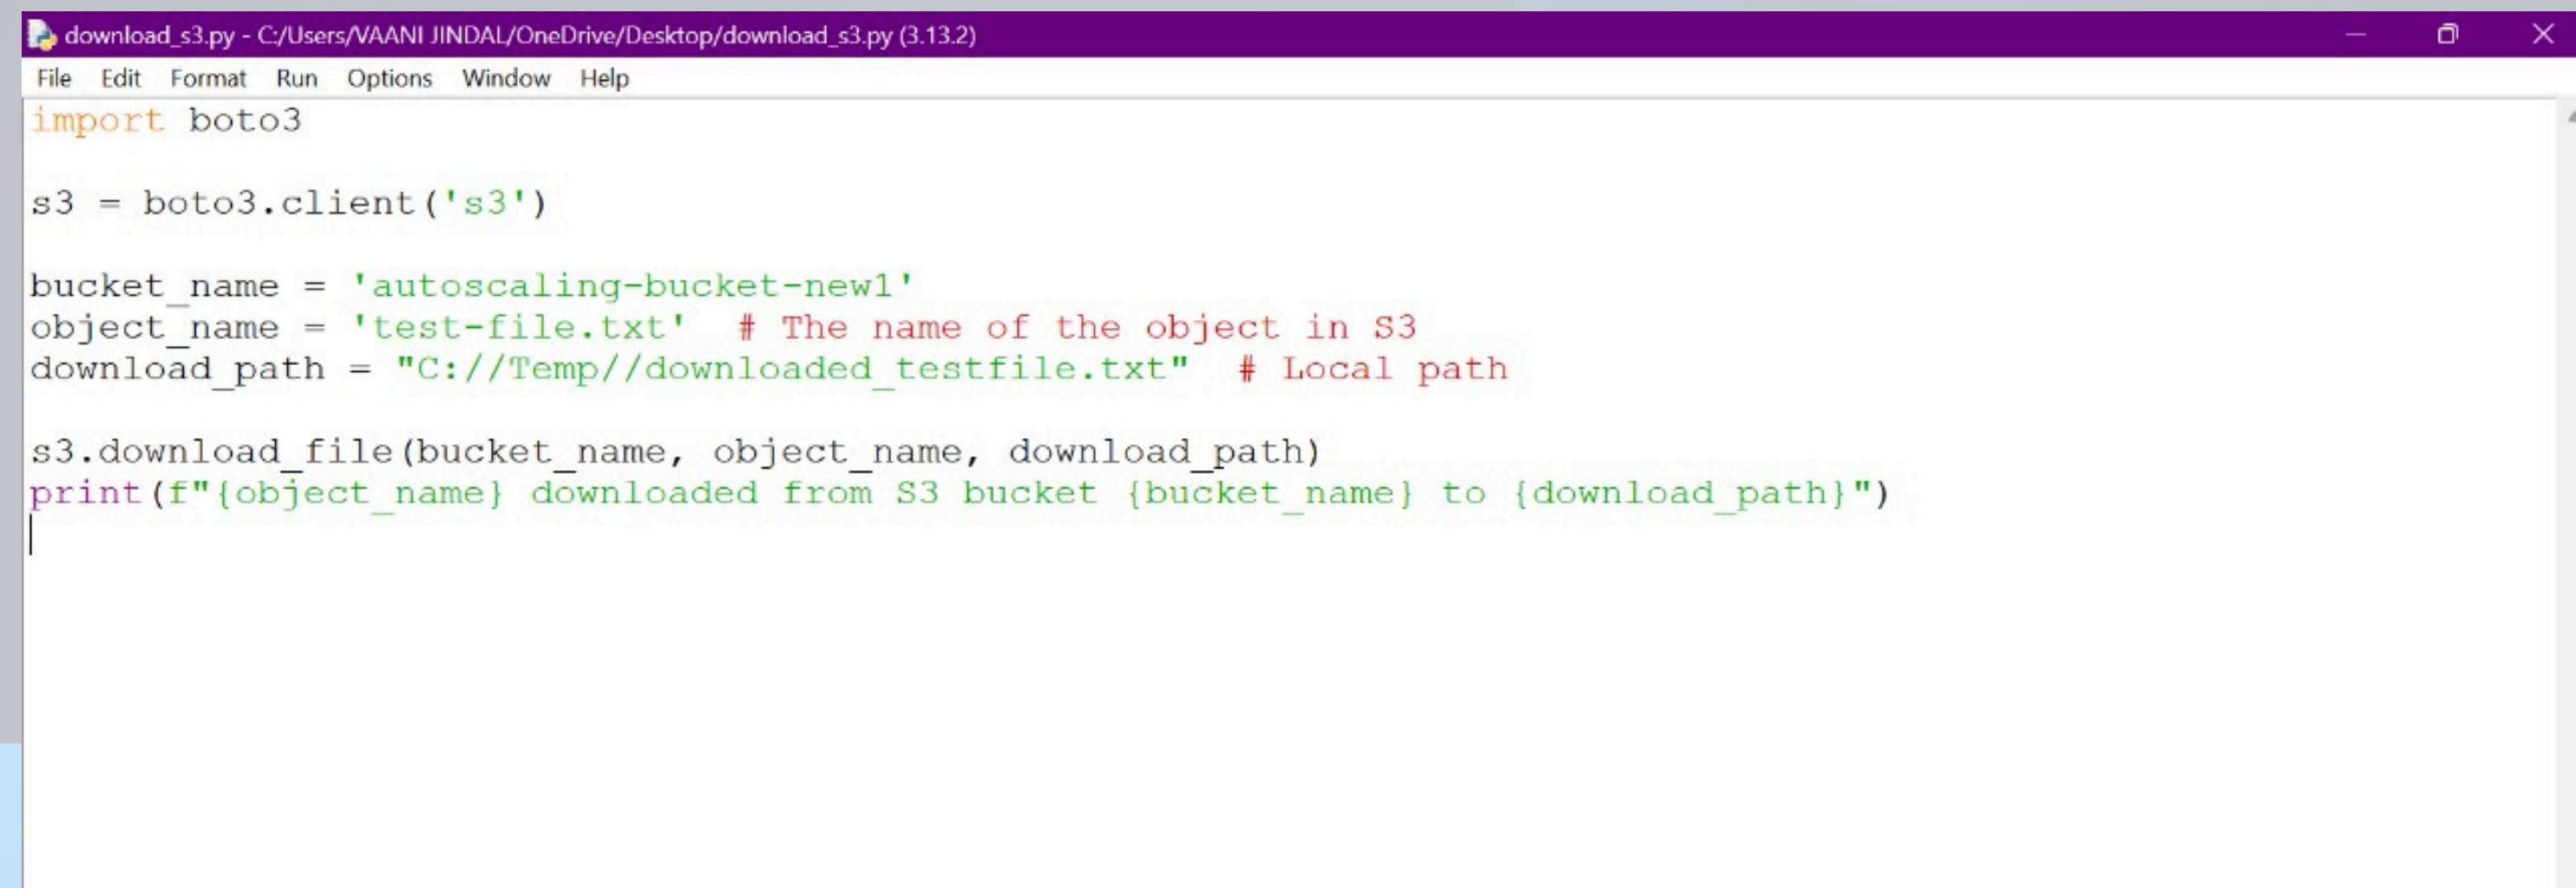


Download a File from S3 using Python

```
C:\Users\VAANI JINDAL\OneDrive\Desktop>mkdir C:\Temp  
  
C:\Users\VAANI JINDAL\OneDrive\Desktop>mkdir C:\Temp  
A subdirectory or file C:\Temp already exists.  
  
C:\Users\VAANI JINDAL\OneDrive\Desktop>python download_s3.py  
test-file.txt downloaded from S3 bucket autoscaling-bucket-new1 to C://Temp//downloaded_testf  
ile.txt
```

Made a directory where to download the file

Download a File from S3 using Python



```
download_s3.py - C:/Users/VAANI JINDAL/OneDrive/Desktop/download_s3.py (3.13.2)
File Edit Format Run Options Window Help
import boto3

s3 = boto3.client('s3')

bucket_name = 'autoscaling-bucket-new1'
object_name = 'test-file.txt' # The name of the object in S3
download_path = "C://Temp//downloaded_testfile.txt" # Local path

s3.download_file(bucket_name, object_name, download_path)
print(f"{object_name} downloaded from S3 bucket {bucket_name} to {download_path}")
|
```

Enable S3 Bucket Versioning or Lifecycle Policies

To ensure data retention and automatic deletion of old files, you should enable S3 versioning and lifecycle policies.

Enabling S3 Bucket Versioning

```
C:\Users\VAANI JINDAL\OneDrive\Desktop>aws s3api put-bucket-versioning --bucket autoscaling-b  
ucket-new1 --versioning-configuration Status=Enabled  
  
C:\Users\VAANI JINDAL\OneDrive\Desktop>aws s3api get-bucket-versioning --bucket autoscaling-b  
ucket-new1  
{  
    "Status": "Enabled"  
}
```

Lifecycle Policies

```
C:\Users\VAANI JINDAL\OneDrive\Desktop>aws s3api put-bucket-lifecycle-configuration --bucket autoscaling-bucket-new1 --lifecycle-configuration file://lifecycle_policy.json
{
  "TransitionDefaultMinimumObjectSize": "all_storage_classes_128K"
}

C:\Users\VAANI JINDAL\OneDrive\Desktop>aws s3api get-bucket-lifecycle-configuration --bucket autoscaling-bucket-new1
{
  "TransitionDefaultMinimumObjectSize": "all_storage_classes_128K",
  "Rules": [
    {
      "Expiration": {
        "Days": 30
      },
      "ID": "DeleteOldFiles",
      "Filter": {
        "Prefix": ""
      },
      "Status": "Enabled"
    }
  ]
}
```

Make
lifecycle_policy
.json file

Command 1: apply the fixed policy command

Command 2: verify the lifecycle policy command

SETTING UP AUTO SCALING FOR EC2 INSTANCES

Launch Template

```
C:\Users\Vanshika Munjal\Desktop>aws ec2 create-launch-template --launch-template-name MyAutoScaleTemplate --version-description "v1"
--launch-template-data "{\"ImageId\": \"ami-0019610b50e1fe6bb\", \"InstanceType\": \"t2.micro\", \"KeyName\": \"MyNewKeypair\", \"SecurityGroupIds\": [\"sg-0d512e764a6a47df0\"], \"IamInstanceProfile\": {\"Name\": \"EC2AutoScalingProfile\"}}"
{
  "LaunchTemplate": {
    "LaunchTemplateId": "lt-08bc6716bc23b51e4",
    "LaunchTemplateName": "MyAutoScaleTemplate",
    "CreateTime": "2025-05-06T21:48:27+00:00",
    "CreatedBy": "arn:aws:iam::426285640896:user/user-autoscaling",
    "DefaultVersionNumber": 1,
    "LatestVersionNumber": 1,
    "Operator": {
      "Managed": false
    }
  }
}
```

```
C:\Users\Vanshika Munjal\Desktop>aws ec2 describe-launch-templates --region ap-south-1
{
  "LaunchTemplates": [
    {
      "LaunchTemplateId": "lt-08bc6716bc23b51e4",
      "LaunchTemplateName": "MyAutoScaleTemplate",
      "CreateTime": "2025-05-06T21:48:27+00:00",
      "CreatedBy": "arn:aws:iam::426285640896:user/user-autoscaling",
      "DefaultVersionNumber": 1,
      "LatestVersionNumber": 1,
      "Operator": {
        "Managed": false
      }
    }
  ]
}
```

This will show all the launch templates including their names and versions

Auto Scaling Group Creation

```
C:\Users\Vanshika Munjal\Desktop>aws autoscaling create-auto-scaling-group --auto-scaling-group-name MyASG --launch-template "LaunchTemplateArn=MyAutoScaleTemplate,Version=1" --min-size 1 --max-size 5 --desired-capacity 2 --vpc-zone-identifier "subnet-09aa416d352d8792e, subnet-081ae8c6100435c9a, subnet-0e68c9beb9282a1ff" --region ap-south-1
```

```
C:\Users\Vanshika Munjal\Desktop>aws autoscaling describe-auto-scaling-groups --auto-scaling-group-names MyASG --region ap-south-1
{
    "AutoScalingGroups": [
        {
            "AutoScalingGroupName": "MyASG",
            "AutoScalingGroupARN": "arn:aws:autoscaling:ap-south-1:426285640896:autoScalingGroup:9064489b-98c8-441c-b80c-2a979972805d:autoScalingGroupName/MyASG",
            "LaunchTemplate": {
                "LaunchTemplateId": "lt-08bc6716bc23b51e4",
                "LaunchTemplateName": "MyAutoScaleTemplate",
                "Version": "1"
            },
            "MinSize": 1,
            "MaxSize": 5,
            "DesiredCapacity": 2,
            "DefaultCooldown": 300,
            "AvailabilityZones": [
                "ap-south-1c",
                "ap-south-1b",
                "ap-south-1a"
            ],
            "LoadBalancerNames": [],
            "TargetGroupARNs": [],
            "HealthCheckType": "EC2",
            "HealthCheckGracePeriod": 0,
            "Instances": [
                {
                    "InstanceId": "i-013429a4d06b536d4",
                    "InstanceType": "t2.micro",
                    "AvailabilityZone": "ap-south-1a",
                    "LifecycleState": "InService",
                    "HealthStatus": "Healthy",
                    "LaunchTemplate": {
                        "LaunchTemplateId": "lt-08bc6716bc23b51e4",
                        "Version": "1"
                    }
                }
            ]
        }
    ]
}
```

Shows the detail of the auto scaling group, allows you to verify that everything is setup correctly.

The screenshot shows the AWS CloudWatch Metrics Insights interface. At the top, there are navigation icons and a search bar labeled "[Alt+S]". Below the search bar, the region is set to "Asia Pacific (Mumbai)" and the user is "Vanshika Munjal". The main area displays a table titled "Instances (3) Info" with the following columns: Name, Instance ID, Instance state, Instance type, Status check, and Alarm status. The table lists three EC2 instances:

Name	Instance ID	Instance state	Instance type	Status check	Alarm status
	i-0283a84daf50a9f3c	Running	t2.micro	2/2 checks passed	View alarms +
	i-013429a4d06b536d4	Running	t2.micro	2/2 checks passed	View alarms +
	i-0361d48e369ecfac5	Running	t2.micro	2/2 checks passed	View alarms +

Minimum number of EC2 instances to keep running at all times.
--max-size 5 Maximum number of EC2 instances that can be launched in the group.
--desired-capacity 2 Number of EC2 instances to launch initially. The Auto Scaling Group will try to maintain this number unless scaling policies change it.

```
C:\Users\Vanshika Munjal\Desktop>aws ec2 describe-vpcs --region ap-south-1
{
    "Vpcs": [
        {
            "OwnerId": "426285640896",
            "InstanceTenancy": "default",
            "CidrBlockAssociationSet": [
                {
                    "AssociationId": "vpc-cidr-assoc-0be3d6771b3dd6738",
                    "CidrBlock": "172.31.0.0/16",
                    "CidrBlockState": {
                        "State": "associated"
                    }
                }
            ],
            "IsDefault": true,
            "BlockPublicAccessStates": {
                "InternetGatewayBlockMode": "off"
            },
            "VpcId": "vpc-08a94d26caf441f1f",
            "State": "available",
            "CidrBlock": "172.31.0.0/16",
            "DhcpOptionsId": "dopt-03b17e549da394ca1"
        }
    ]
}
```

This command will return a list of all VPCs in ap-south-1 region including their IDs

Create a Target Group

```
C:\Users\Vanshika Munjal\Desktop>aws elbv2 create-target-group --name MyTargetGroup --protocol HTTP --port 80 --vpc-id vpc-08a94d26ca  
f441f1f --region ap-south-1  
{  
    "TargetGroups": [  
        {  
            "TargetGroupArn": "arn:aws:elasticloadbalancing:ap-south-1:426285640896:targetgroup/MyTargetGroup/6df82defd1ccdb7e",  
            "TargetGroupName": "MyTargetGroup",  
            "Protocol": "HTTP",  
            "Port": 80,  
            "VpcId": "vpc-08a94d26caf441f1f",  
            "HealthCheckProtocol": "HTTP",  
            "HealthCheckPort": "traffic-port",  
            "HealthCheckEnabled": true,  
            "HealthCheckIntervalSeconds": 30,  
            "HealthCheckTimeoutSeconds": 5,  
            "HealthyThresholdCount": 5,  
            "UnhealthyThresholdCount": 2,  
            "HealthCheckPath": "/",  
            "Matcher": {  
                "HttpCode": "200"  
            },  
            "TargetType": "instance",  
            "ProtocolVersion": "HTTP1",  
            "IpAddressType": "ipv4"  
        }  
    ]  
}
```

LOAD BALANCER

Attaching Load Balancer

```
C:\Users\Vanshika Munjal\Desktop>aws elbv2 create-load-balancer ^ --name MyAppLoadBalancer ^ --subnets subnet-0e68c9beb9282a1ff subnet-09aa416d352d8792e subnet-081ae8c6100435c9a ^ --security-groups sg-0d512e764a6a47df0 ^ --scheme internet-facing ^ --type application ^ --ip-address-type ipv4 ^ --region ap-south-1
{
    "LoadBalancers": [
        {
            "LoadBalancerArn": "arn:aws:elasticloadbalancing:ap-south-1:426285640896:loadbalancer/app/MyAppLoadBalancer/b4eda167e52513da",
            "DNSName": "MyAppLoadBalancer-1054520609.ap-south-1.elb.amazonaws.com",
            "CanonicalHostedZoneId": "ZP97RAFLXTNZK",
            "CreatedTime": "2025-05-06T22:14:26.780000+00:00",
            "LoadBalancerName": "MyAppLoadBalancer",
            "Scheme": "internet-facing",
            "VpcId": "vpc-08a94d26caf441f1f",
            "State": {
                "Code": "provisioning"
            },
            "Type": "application",
            "AvailabilityZones": [
                {
                    "ZoneName": "ap-south-1c",
                    "SubnetId": "subnet-0e68c9beb9282a1ff",
                    "LoadBalancerAddresses": []
                },
                {
                    "ZoneName": "ap-south-1a",
                    "SubnetId": "subnet-09aa416d352d8792e",
                    "LoadBalancerAddresses": []
                },
                {
                    "ZoneName": "ap-south-1b",
                    "SubnetId": "subnet-081ae8c6100435c9a",
                    "LoadBalancerAddresses": []
                }
            ],
            "SecurityGroups": [
                "sg-0d512e764a6a47df0"
            ]
        }
    ]
}
```

Creating a listener

```
C:\Users\Vanshika Munjal\Desktop>aws elbv2 create-listener ^ --load-balancer-arn arn:aws:elasticloadbalancing:ap-south-1:426285640896:loadbalancer/app/MyAppLoadBalancer/b4eda167e52513da ^ --protocol HTTP ^ --port 80 ^ --default-actions Type=forward,TargetGroupArn=arn:aws:elasticloadbalancing:ap-south-1:426285640896:targetgroup/MyTargetGroup/6df82defd1ccdb7e ^ --region ap-south-1
{
  "Listeners": [
    {
      "ListenerArn": "arn:aws:elasticloadbalancing:ap-south-1:426285640896:listener/app/MyAppLoadBalancer/b4eda167e52513da/4e60ced0cc4911b7",
      "LoadBalancerArn": "arn:aws:elasticloadbalancing:ap-south-1:426285640896:loadbalancer/app/MyAppLoadBalancer/b4eda167e52513da",
      "Port": 80,
      "Protocol": "HTTP",
      "DefaultActions": [
        {
          "Type": "forward",
          "TargetGroupArn": "arn:aws:elasticloadbalancing:ap-south-1:426285640896:targetgroup/MyTargetGroup/6df82defd1ccdb7e",
          "ForwardConfig": {
            "TargetGroups": [
              {
                "TargetGroupArn": "arn:aws:elasticloadbalancing:ap-south-1:426285640896:targetgroup/MyTargetGroup/6df82defd1ccdb7e",
                "Weight": 1
              }
            ],
            "TargetGroupStickinessConfig": {
              "Enabled": false
            }
          }
        }
      ]
    }
  ]
}
```

Registering instances to target group

```
C:\Users\Vanshika Munjal\Desktop>aws elbv2 register-targets --target-group-arn arn:aws:elasticloadbalancing:ap-south-1:426285640896:targetgroup/MyTargetGroup/6df82defd1ccdb7e --targets Id=i-0283a84daf50a9f3c Id=i-013429a4d06b536d4 Id=i-0361d48e369ecfac5 --region ap-south-1
```

Attach Target Group to Auto Scaling Group

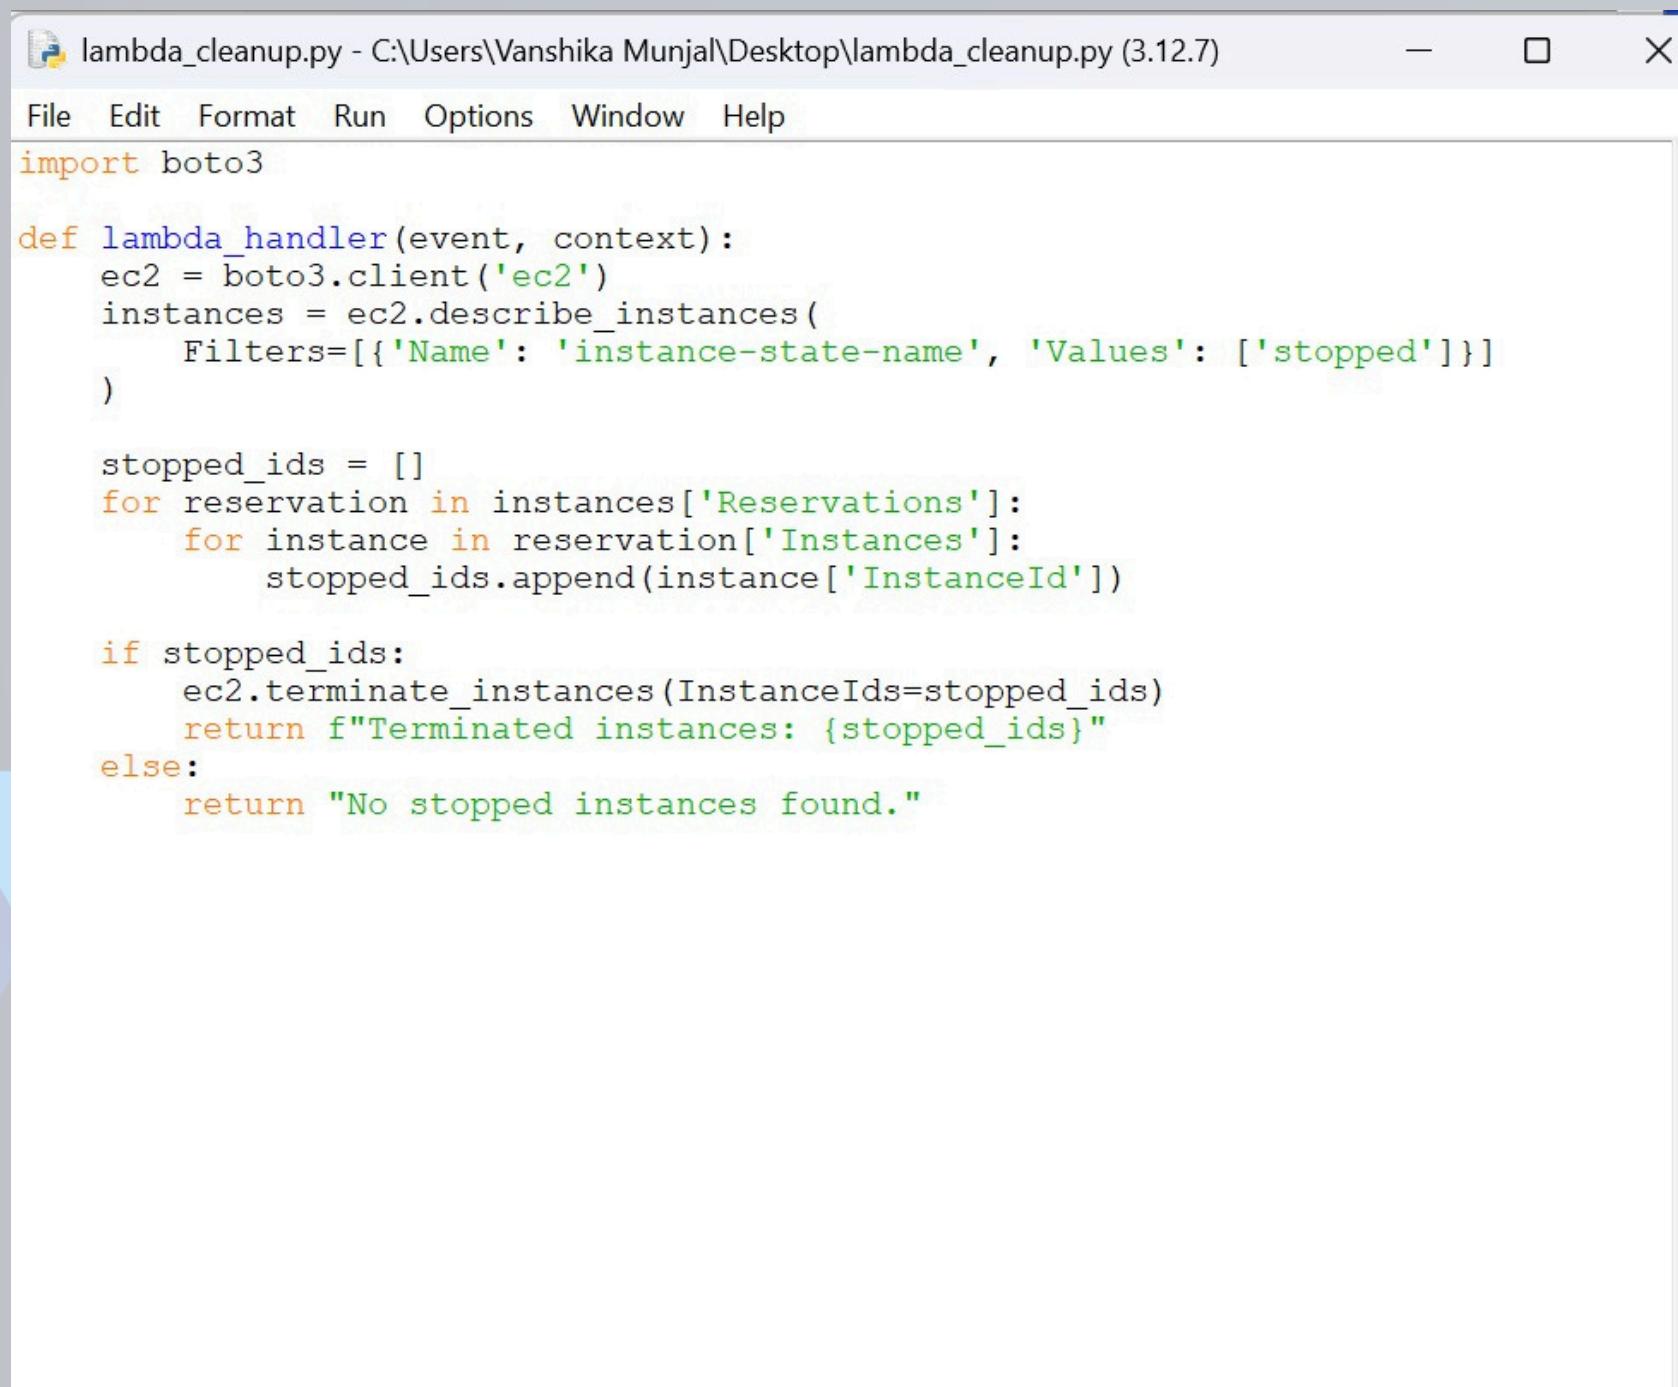
```
C:\Users\Vanshika Munjal\Desktop>aws autoscaling attach-load-balancer-target-groups ^ --auto-scaling-group-name MyASG ^ --target-groups arn:aws:elasticloadbalancing:ap-south-1:426285640896:targetgroup/MyTargetGroup/6df82defd1ccdb7e ^ --region ap-south-1
```

Describing Load Balancer Target Groups

```
C:\Users\Vanshika Munjal\Desktop>aws autoscaling describe-load-balancer-target-groups --auto-scaling-group-name MyASG --region ap-south-1
{
  "LoadBalancerTargetGroups": [
    {
      "LoadBalancerTargetGroupARN": "arn:aws:elasticloadbalancing:ap-south-1:426285640896:targetgroup/MyTargetGroup/6df82defd1ccdb7e",
      "State": "Added"
    }
  ]
}
```

LAMBDA FUNCTION

Python file



The screenshot shows a Python code editor window titled "lambda_cleanup.py - C:\Users\Vanshika Munjal\Desktop\lambda_cleanup.py (3.12.7)". The menu bar includes File, Edit, Format, Run, Options, Window, and Help. The code itself is a Lambda function that uses the boto3 library to interact with EC2. It defines a handler function that filters stopped EC2 instances and terminates them. The code is as follows:

```
lambda_cleanup.py - C:\Users\Vanshika Munjal\Desktop\lambda_cleanup.py (3.12.7)
File Edit Format Run Options Window Help
import boto3

def lambda_handler(event, context):
    ec2 = boto3.client('ec2')
    instances = ec2.describe_instances(
        Filters=[{'Name': 'instance-state-name', 'Values': ['stopped']}]
    )

    stopped_ids = []
    for reservation in instances['Reservations']:
        for instance in reservation['Instances']:
            stopped_ids.append(instance['InstanceId'])

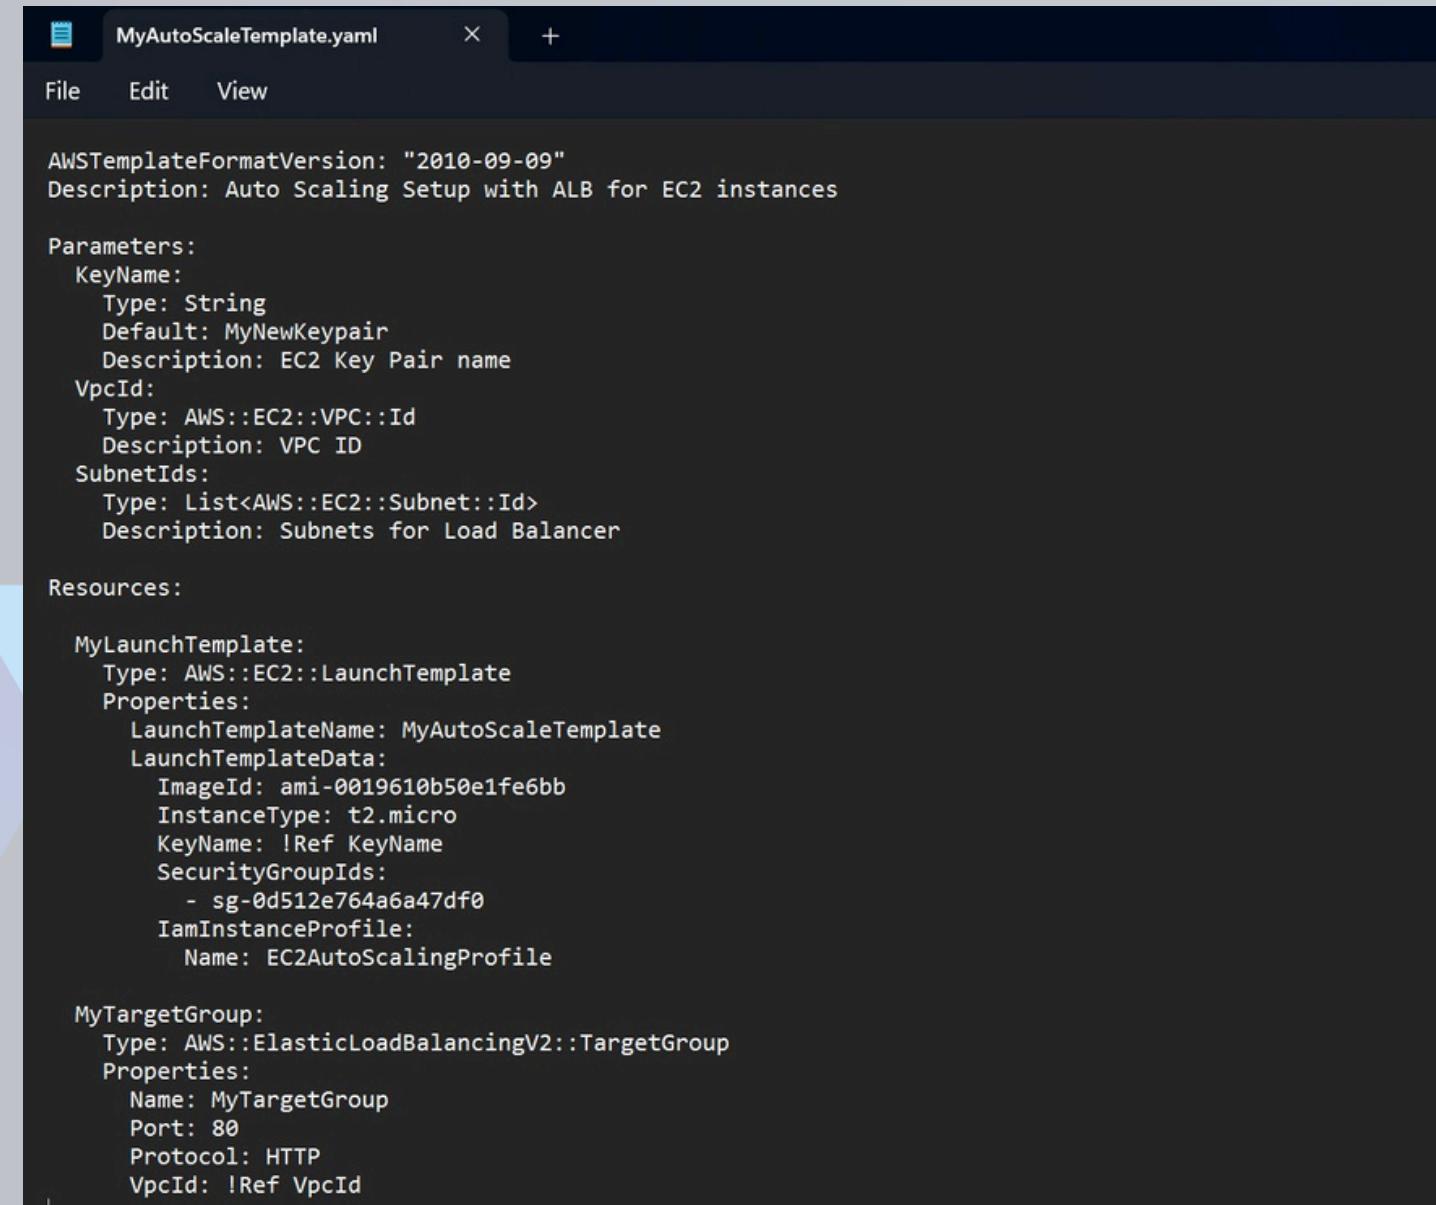
    if stopped_ids:
        ec2.terminate_instances(InstanceIds=stopped_ids)
        return f"Terminated instances: {stopped_ids}"
    else:
        return "No stopped instances found."
```

Lambda Function
deletes stopped EC2
Instances to save cost-
runs on schedule using
CloudWatch Events

Package Lambda Function

```
C:\Users\Vanshika Munjal\Desktop>aws cloudformation package ^ --template-file MyAutoScaleTemplate.yaml ^ --s3-bucket autoscaling-bucket-newl ^ --output-template-file output-template.yaml ^ --region ap-south-1

Successfully packaged artifacts and wrote output template to file output-template.yaml.
Execute the following command to deploy the packaged template
aws cloudformation deploy --template-file C:\Users\Vanshika Munjal\Desktop\output-template.yaml --stack-name <YOUR STACK NAME>
```



The screenshot shows a code editor window titled "MyAutoScaleTemplate.yaml". The file contains a CloudFormation template with the following structure:

```
AWSTemplateFormatVersion: "2010-09-09"
Description: Auto Scaling Setup with ALB for EC2 instances

Parameters:
  KeyName:
    Type: String
    Default: MyNewKeypair
    Description: EC2 Key Pair name
  VpcId:
    Type: AWS::EC2::VPC::Id
    Description: VPC ID
  SubnetIds:
    Type: List<AWS::EC2::Subnet::Id>
    Description: Subnets for Load Balancer

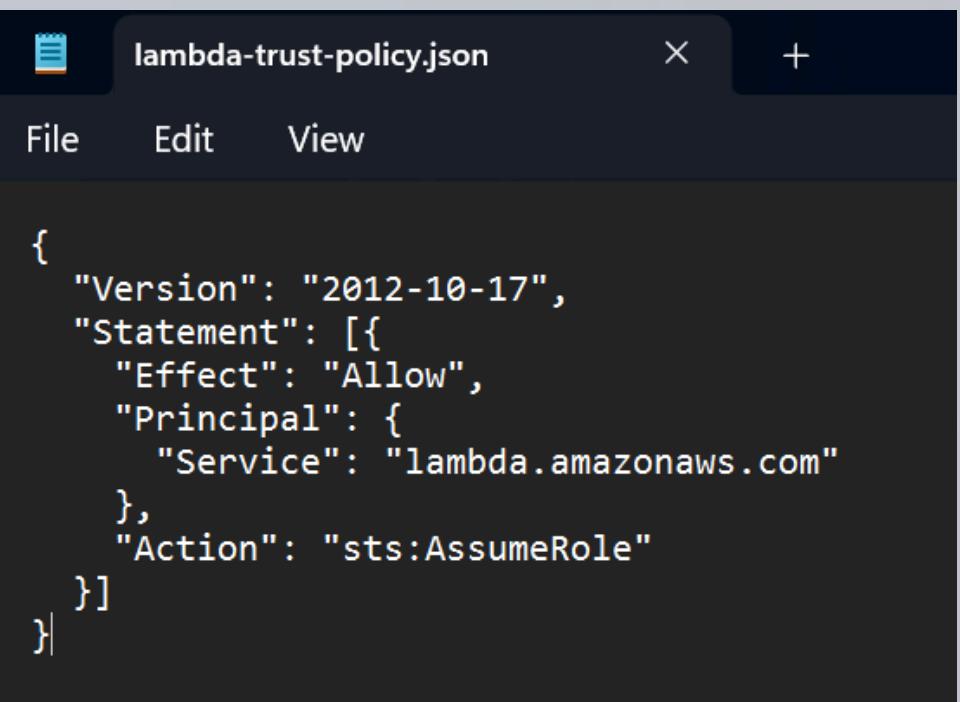
Resources:
  MyLaunchTemplate:
    Type: AWS::EC2::LaunchTemplate
    Properties:
      LaunchTemplateName: MyAutoScaleTemplate
      LaunchTemplateData:
        ImageId: ami-0019610b50e1fe6bb
        InstanceType: t2.micro
        KeyName: !Ref KeyName
        SecurityGroupIds:
          - sg-0d512e764a6a47df0
        IamInstanceProfile:
          Name: EC2AutoScalingProfile

  MyTargetGroup:
    Type: AWS::ElasticLoadBalancingV2::TargetGroup
    Properties:
      Name: MyTargetGroup
      Port: 80
      Protocol: HTTP
      VpcId: !Ref VpcId
```

Defines the template format version and a brief description

```
C:\Users\Vanshika Munjal\Desktop>aws iam create-role --role-name LambdaEC2CleanupRole --assume-role-policy-document file://lambda-trust-policy.json
{
  "Role": {
    "Path": "/",
    "RoleName": "LambdaEC2CleanupRole",
    "RoleId": "AROAWGQE2VTA0IJE55CUY",
    "Arn": "arn:aws:iam::426285640896:role/LambdaEC2CleanupRole",
    "CreateDate": "2025-05-07T11:31:50+00:00",
    "AssumeRolePolicyDocument": {
      "Version": "2012-10-17",
      "Statement": [
        {
          "Effect": "Allow",
          "Principal": {
            "Service": "lambda.amazonaws.com"
          },
          "Action": "sts:AssumeRole"
        }
      ]
    }
  }
}
```

Creating IAM Role for Lambda Function



A screenshot of a code editor window titled "lambda-trust-policy.json". The window shows a JSON policy document with the following content:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "Service": "lambda.amazonaws.com"
      },
      "Action": "sts:AssumeRole"
    }
  ]
}
```

Attach Policies

```
C:\Users\Vanshika Munjal\Desktop>aws iam attach-role-policy --role-name LambdaEC2CleanupRole --policy-arn arn:aws:iam::aws:policy/AmazonEC2FullAccess
```

```
C:\Users\Vanshika Munjal\Desktop>aws lambda create-function --function-name EC2CleanupFunction --runtime python3.12 --role arn:iam::426285640896:role/LambdaEC2CleanupRole --handler lambda_cleanup.lambda_handler --zip-file fileb://lambda_cleanup.zip --region ap-south-1

{
  "FunctionName": "EC2CleanupFunction",
  "FunctionArn": "arn:aws:lambda:ap-south-1:426285640896:function:EC2CleanupFunction",
  "Runtime": "python3.12",
  "Role": "arn:aws:iam::426285640896:role/LambdaEC2CleanupRole",
  "Handler": "lambda_cleanup.lambda_handler",
  "CodeSize": 487,
  "Description": "",
  "Timeout": 3,
  "MemorySize": 128,
  "LastModified": "2025-05-07T11:36:38.790+0000",
  "CodeSha256": "fG6q7Ca8ap/kjrJ3kdSQloIquistiQoeXqwF4JLJz+Zo=",
  "Version": "$LATEST",
  "TracingConfig": {
    "Mode": "PassThrough"
  },
  "RevisionId": "2ecdaebc-425a-4487-b924-f6f0c7efcb77",
  "State": "Pending",
  "StateReason": "The function is being created.",
  "StateReasonCode": "Creating",
  "PackageType": "Zip",
  "Architectures": [
    "x86_64"
  ],
  "EphemeralStorage": {
    "Size": 512
  },
  "SnapStart": {
    "ApplyOn": "None",
    "OptimizationStatus": "Off"
  }
},
```

Deploy Lambda Function

Schedule Lambda with Cloud Watch Events

```
C:\Users\Vanshika Munjal\Desktop>aws events put-rule --schedule-expression "rate(1 hour)" --name EC2CleanupRule
{
  "RuleArn": "arn:aws:events:ap-south-1:426285640896:rule/EC2CleanupRule"
}
```

Add Lambda as Target

```
C:\Users\Vanshika Munjal\Desktop>aws events put-targets ^ --rule EC2CleanupRule ^ --targets Id="1",Arn="arn:aws:lambda:ap-south-1:426285640896:function:EC2CleanupFunction"
{
    "FailedEntryCount": 0,
    "FailedEntries": []
}
```

Give permissions to event bridge

```
C:\Users\Vanshika Munjal\Desktop>aws lambda add-permission ^ --function-name EC2CleanupFunction ^ --statement-id AllowExecutionFromEventBridge ^ --action lambda:InvokeFunction ^ --principal events.amazonaws.com ^ --source-arn arn:aws:events:ap-south-1:426285640896:rule/EC2CleanupRule ^ --region ap-south-1
{
    "Statement": "{\"Sid\":\"AllowExecutionFromEventBridge\",\"Effect\":\"Allow\",\"Principal\":{\"Service\":\"events.amazonaws.com\"},\"Action\":\"lambda:InvokeFunction\",\"Resource\":\"arn:aws:lambda:ap-south-1:426285640896:function:EC2CleanupFunction\",\"Condition\":{\"ArnLike\":{\"AWS:SourceArn\":\"arn:aws:events:ap-south-1:426285640896:rule/EC2CleanupRule\"}}}"
}
```

The screenshot shows the AWS Lambda Functions interface. At the top, the navigation bar includes 'Lambda' > 'Functions' > 'EC2CleanupFunction'. The main title is 'EC2CleanupFunction'. Below it, the 'Function overview' section is expanded, showing a diagram where 'EC2CleanupFunction' is triggered by 'EventBridge (CloudWatch Events)'. The function has no layers. A 'Diagram' tab is selected, and a 'Template' tab is available. On the right, there are tabs for 'Info' (selected) and 'Tutorials'. The 'Info' tab displays the function's ARN: arn:aws:lambda:ap-south-1:426285640896:function:EC2CleanupFunction. The 'Tutorials' tab shows a section titled 'Create a simple web app' with a description: 'Learn how to implement common use cases in AWS Lambda.' It lists steps to build a simple web app using Lambda and provides a 'Start tutorial' button.

You now have a Lambda function running hourly, scanning for unused EC2 instances (such as those spun up by Auto Scaling and later stopped), and terminating them automatically.

CHALLENGES FACED & SOLUTIONS IMPLEMENTED

1. IAM Role Permissions

Challenge: Lambda functions and EC2 instances initially failed to access required AWS resources (S3, CloudWatch).

Solution: Created and attached custom IAM roles with the minimum required permissions. Verified using the IAM Policy Simulator.

2. Lambda Permission Error

Challenge: EventBridge rule couldn't trigger the Lambda function.

Solution: Used aws lambda add-permission to explicitly allow EventBridge to invoke the Lambda function with the correct source-arn.

Solution: Rechecked the full ARN in the Lambda console and replaced placeholders with the actual region and account ID.

CHALLENGES FACED & SOLUTIONS IMPLEMENTED

3. Invalid Target Group ARN

Challenge: Encountered ARN format errors while registering the Lambda function to EventBridge.

Solution: Rechecked the full ARN in the Lambda console and replaced placeholders with the actual region and account ID.

4. Load Balancer & Subnet Mismatch

Challenge: ELB failed to associate with the Auto Scaling Group due to incorrect or insufficient subnets.

Solution: Used aws ec2 describe-subnets to confirm subnets were in different Availability Zones and belonged to the correct VPC.

FUTURE IMPROVEMENTS

- Implement Scaling Policies: Add CPU or memory-based scaling policies to make Auto Scaling more intelligent.
- Add CI/CD Integration: Use AWS CodePipeline and CodeDeploy to automate deployments for Lambda and EC2.
- Introduce Step Scaling or Scheduled Scaling: Optimize resource usage during known traffic patterns.
- Implement HTTPS on ELB: Add SSL certificates to support secure connections.
- Enable Detailed Monitoring: Use detailed metrics in CloudWatch for better insights and alarms.

CONCLUSION

This project successfully demonstrated a dynamic and scalable cloud architecture using AWS services such as EC2, Auto Scaling, Elastic Load Balancer, IAM, S3, EBS, Lambda, and CloudWatch. The infrastructure auto-scales based on demand, ensures high availability, and is integrated with automation through EventBridge and Lambda functions.

THANK YOU